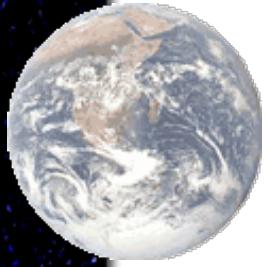


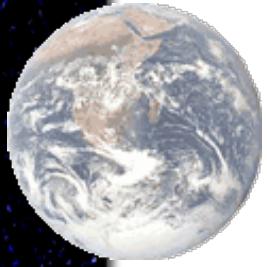
File Systems



File Systems

Essential requirements for long-term information storage:

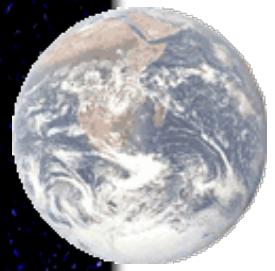
- It must be possible to store a very large amount of information.
- The information must survive the termination of the process using it.
- Multiple processes must be able to access the information concurrently.



File Systems

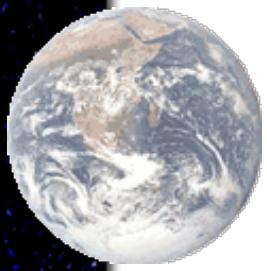
Think of a disk as a linear sequence of fixed-size blocks and supporting reading and writing of blocks. Questions that quickly arise:

- How do you find information?
- How do you keep one user from reading another's data?
- How do you know which blocks are free?

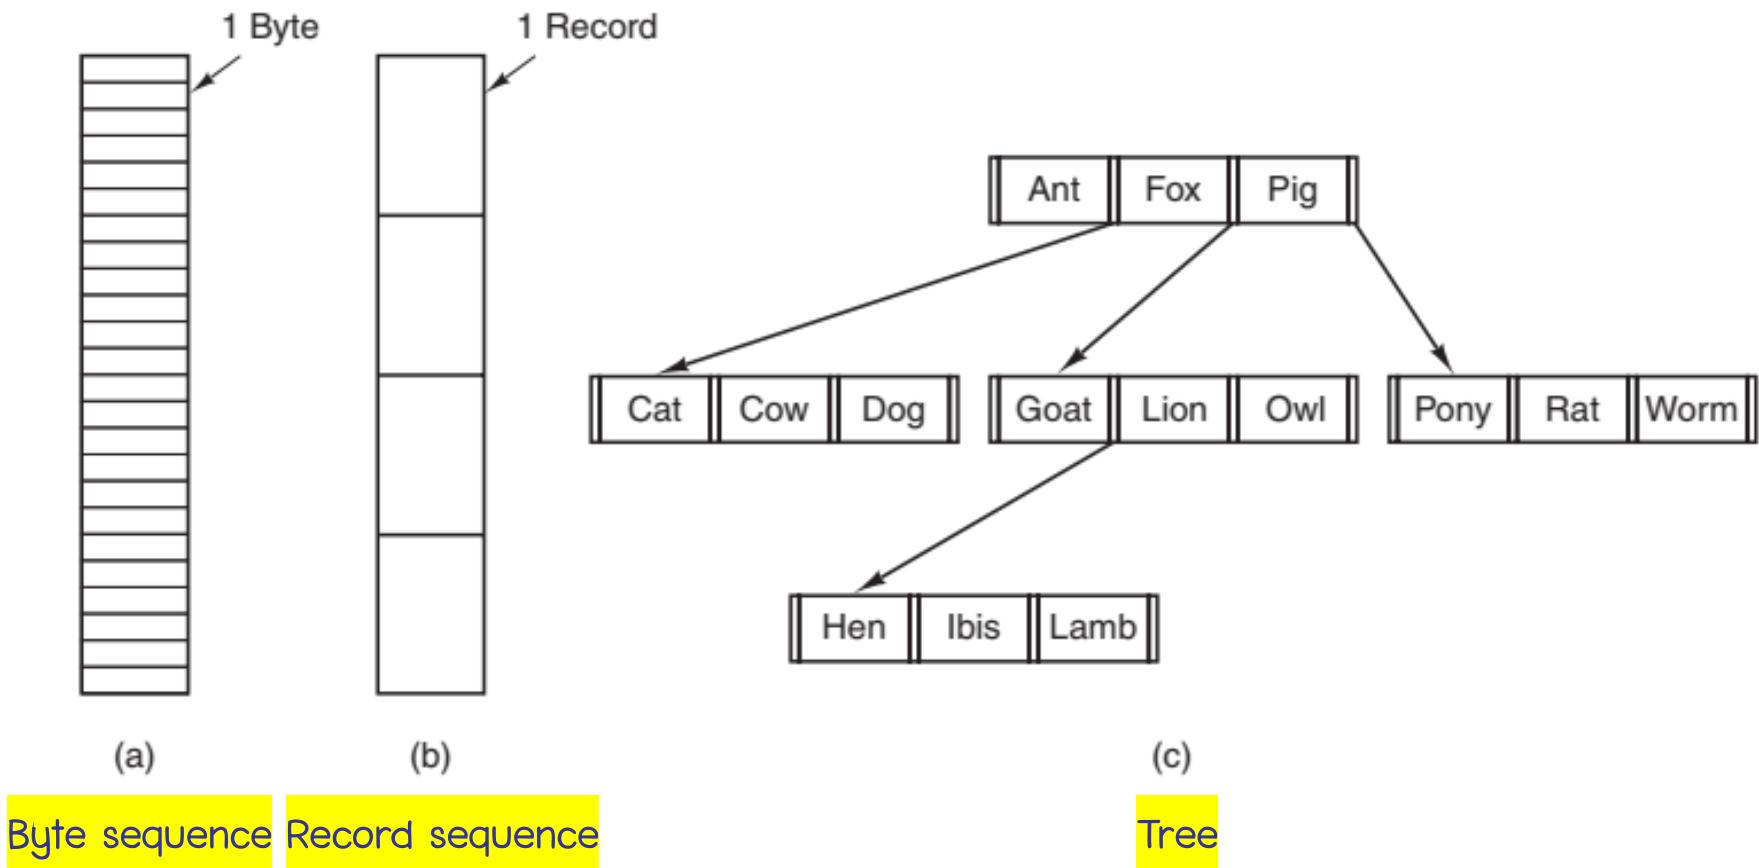


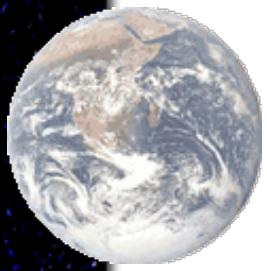
Typical File Extensions

Extension	Meaning
.bak	Backup file
.c	C source program
.gif	Compuserve Graphical Interchange Format image
.hlp	Help file
.html	World Wide Web HyperText Markup Language document
.jpg	Still picture encoded with the JPEG standard
.mp3	Music encoded in MPEG layer 3 audio format
.mpg	Movie encoded with the MPEG standard
.o	Object file (compiler output, not yet linked)
.pdf	Portable Document Format file
.ps	PostScript file
.tex	Input for the TEX formatting program
.txt	General text file
.zip	Compressed archive

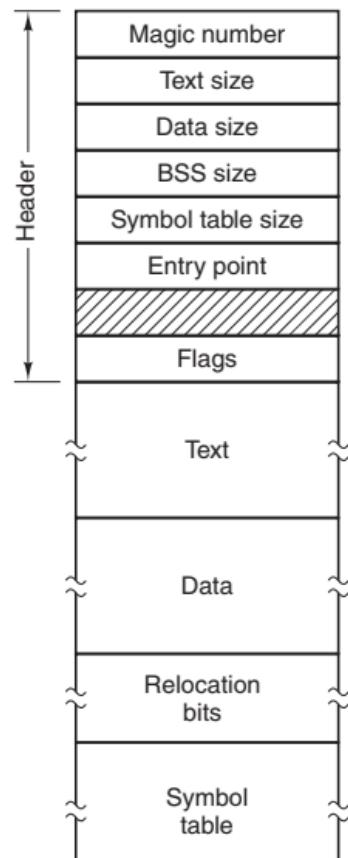


File Structure



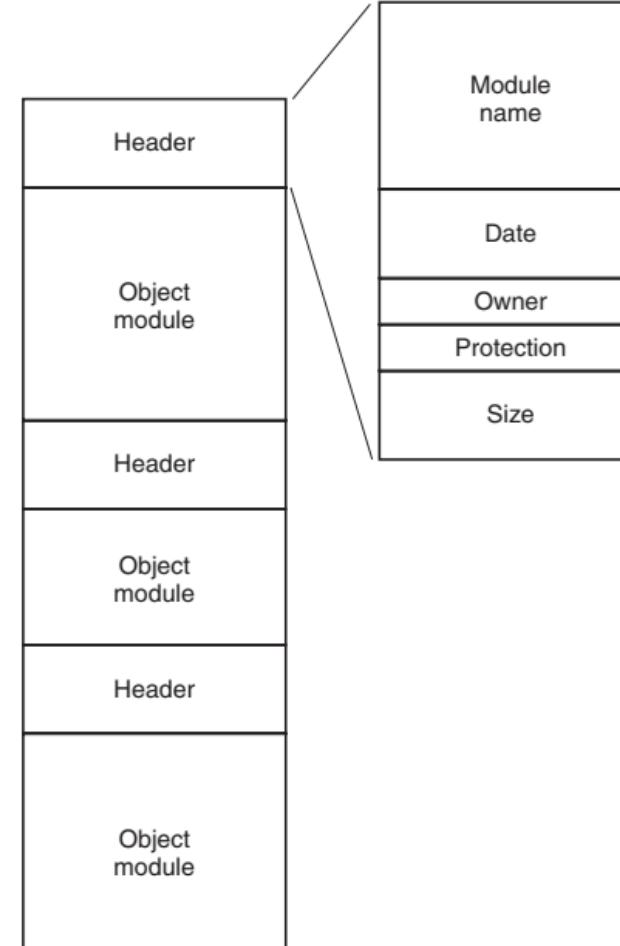


File Types



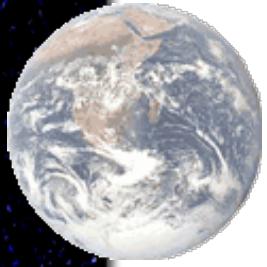
(a)

An executable file



(b)

An archive



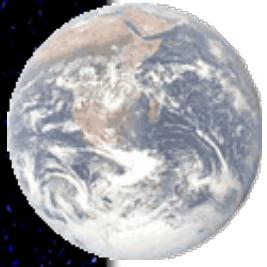
Some Possible File Attributes

Attribute	Meaning
Protection	Who can access the file and in what way
Password	Password needed to access the file
Creator	ID of the person who created the file
Owner	Current owner
Read-only flag	0 for read/write; 1 for read only
Hidden flag	0 for normal; 1 for do not display in listings
System flag	0 for normal files; 1 for system file
Archive flag	0 for has been backed up; 1 for needs to be backed up
ASCII/binary flag	0 for ASCII file; 1 for binary file
Random access flag	0 for sequential access only; 1 for random access
Temporary flag	0 for normal; 1 for delete file on process exit
Lock flags	0 for unlocked; nonzero for locked
Record length	Number of bytes in a record
Key position	Offset of the key within each record
Key length	Number of bytes in the key field
Creation time	Date and time the file was created
Time of last access	Date and time the file was last accessed
Time of last change	Date and time the file was last changed
Current size	Number of bytes in the file
Maximum size	Number of bytes the file may grow to



File Operations

- The most common system calls relating to files:
 - Create
 - Delete
 - Open
 - Close
 - Read
 - Write
 - Append
 - Seek
 - Get Attributes
 - Set Attributes
 - Rename



Example Program Using File System Calls

```
/* File copy program. Error checking and reporting is minimal. */

#include <sys/types.h>                                /* include necessary header files */
#include <fcntl.h>
#include <stdlib.h>
#include <unistd.h>

int main(int argc, char *argv[]);                      /* ANSI prototype */

#define BUF_SIZE 4096                                     /* use a buffer size of 4096 bytes */
#define OUTPUT_MODE 0700                                  /* protection bits for output file */

int main(int argc, char *argv[])
{
    int in_fd, out_fd, rd_count, wt_count;
    char buffer[BUF_SIZE];

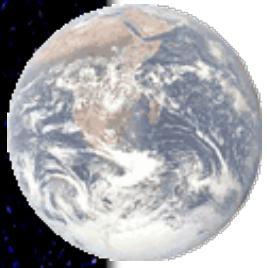
    if (argc != 3) exit(1);                               /* syntax error if argc is not 3 */

    /* Open the input file and create the output file */
    in_fd = open(argv[1], O_RDONLY);                     /* open the source file */
    if (in_fd < 0) exit(2);                             /* if it cannot be opened, exit */
    out_fd = creat(argv[2], OUTPUT_MODE);                /* create the destination file */
    if (out_fd < 0) exit(3);                            /* if it cannot be created, exit */

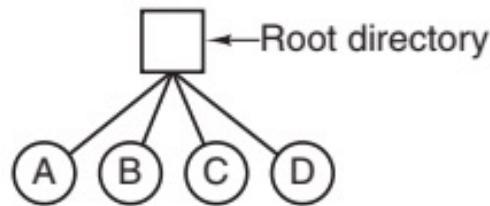
    /* Copy loop */
    while (TRUE) {
        rd_count = read(in_fd, buffer, BUF_SIZE); /* read a block of data */
        if (rd_count <= 0) break;                  /* if end of file or error, exit loop */
        wt_count = write(out_fd, buffer, rd_count); /* write data */
        if (wt_count <= 0) exit(4);                /* wt_count <= 0 is an error */
    }

    /* Close the files */
    close(in_fd);
    close(out_fd);
    if (rd_count == 0)                                /* no error on last read */
        exit(0);
    else
        exit(5);                                    /* error on last read */
}
```

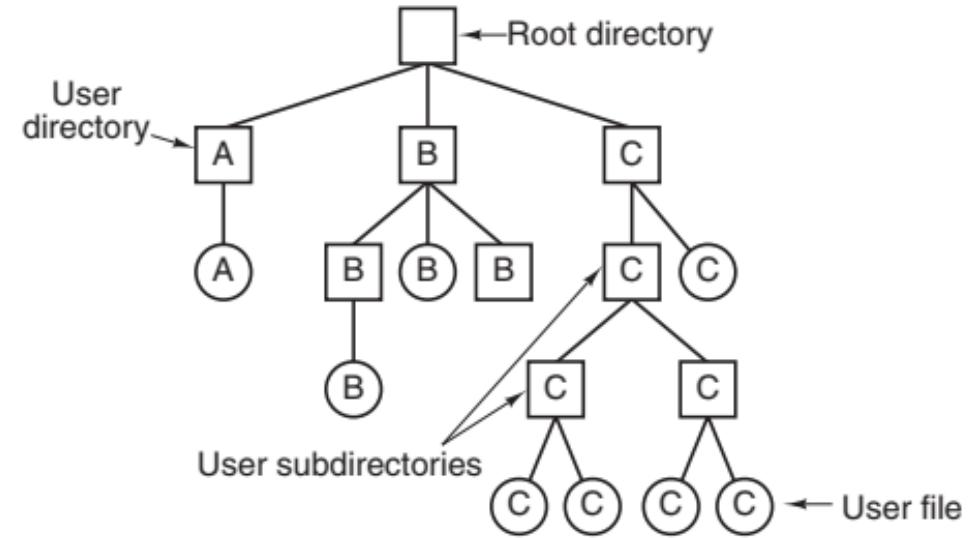
Program copyfile:
copyfile fileA fileB



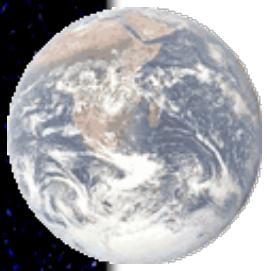
Directory Systems



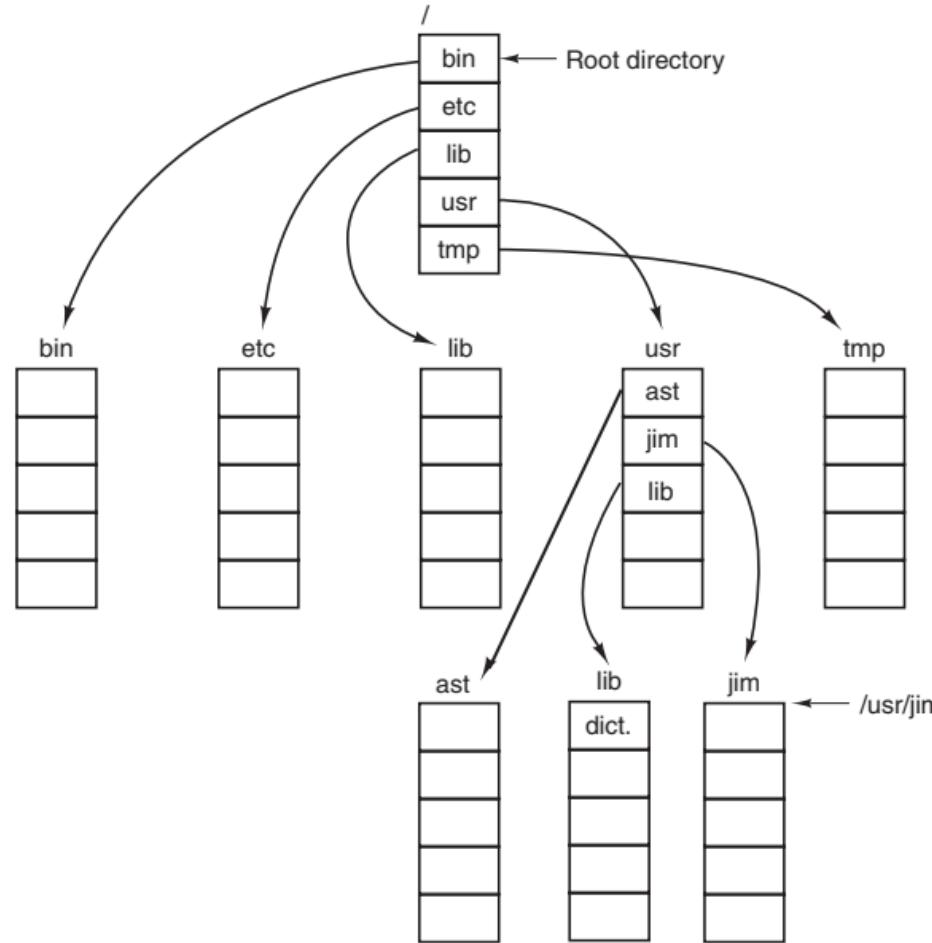
A single-level directory system



A hierarchical directory system



Path Names



A UNIX directory tree



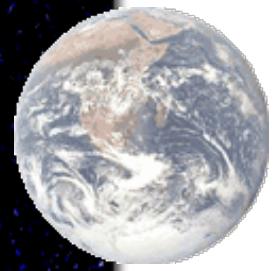
Directory Operations

- System calls for managing directories:
 - Create
 - Delete
 - Opendir
 - Closedir
 - Readdir
 - Rename
 - Link
 - Unlink

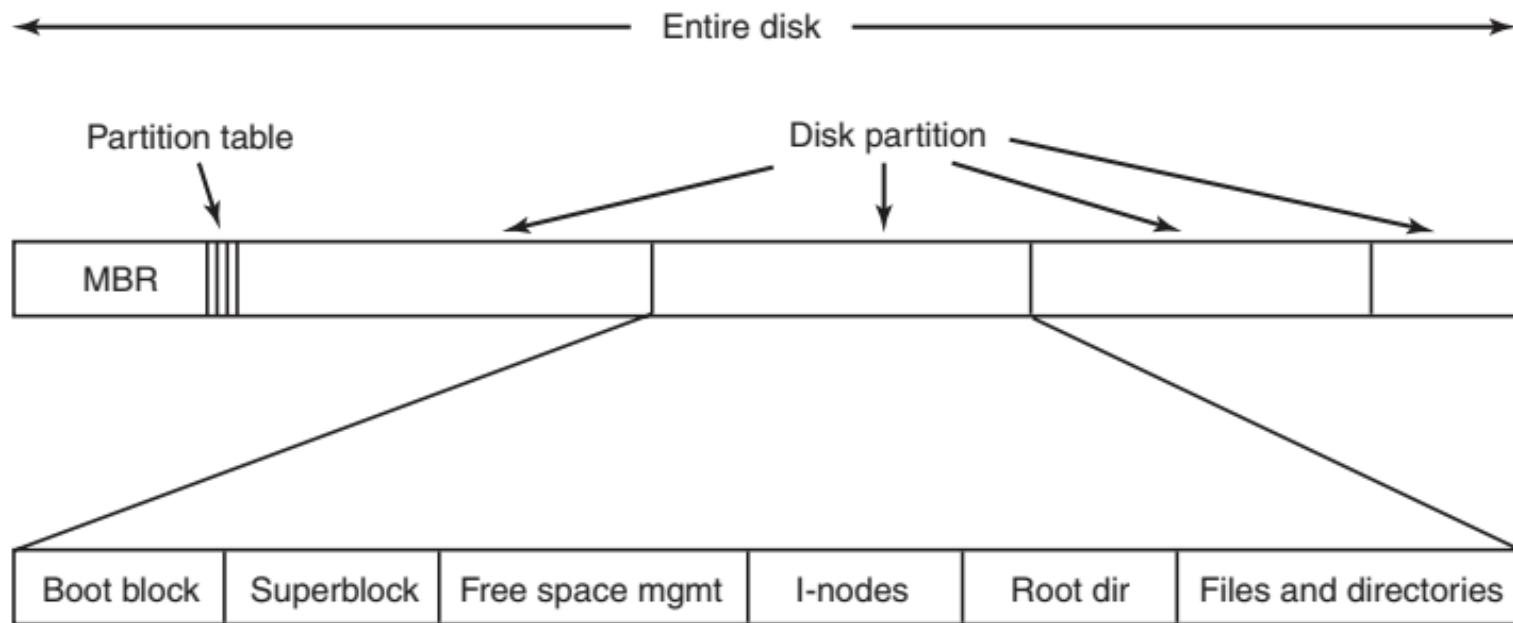


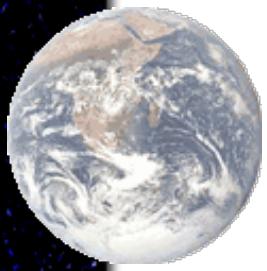
File System Implementation

- File System Layout
- Implementing File Systems
 - Contiguous allocation
 - Linked-list allocation
 - Linked-list allocation using table in memory
 - I-nodes
- Implementing Directories
- Shared Files

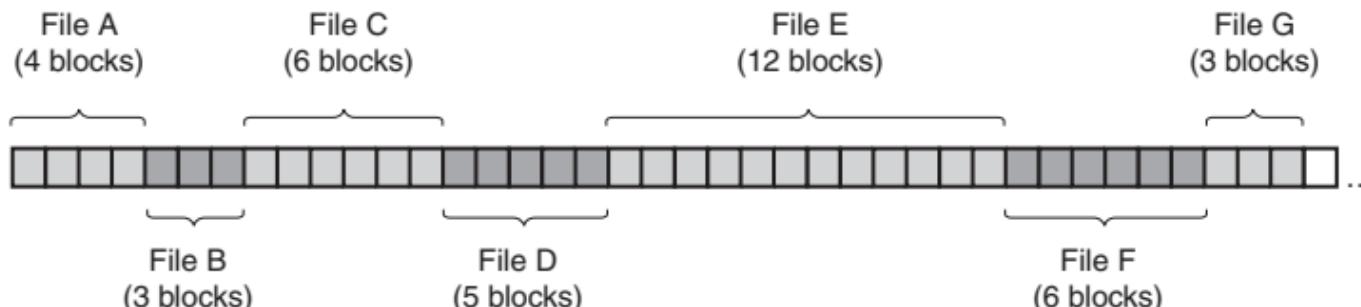


A Possible File System Layout

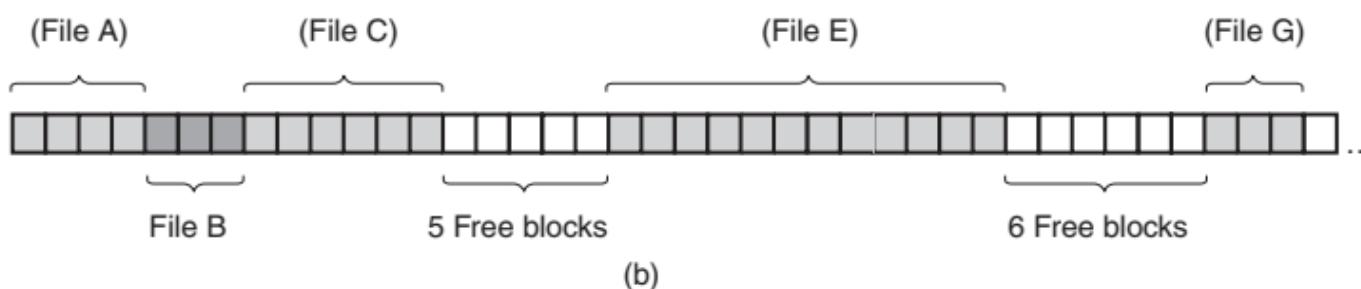




Contiguous Allocation



(a)



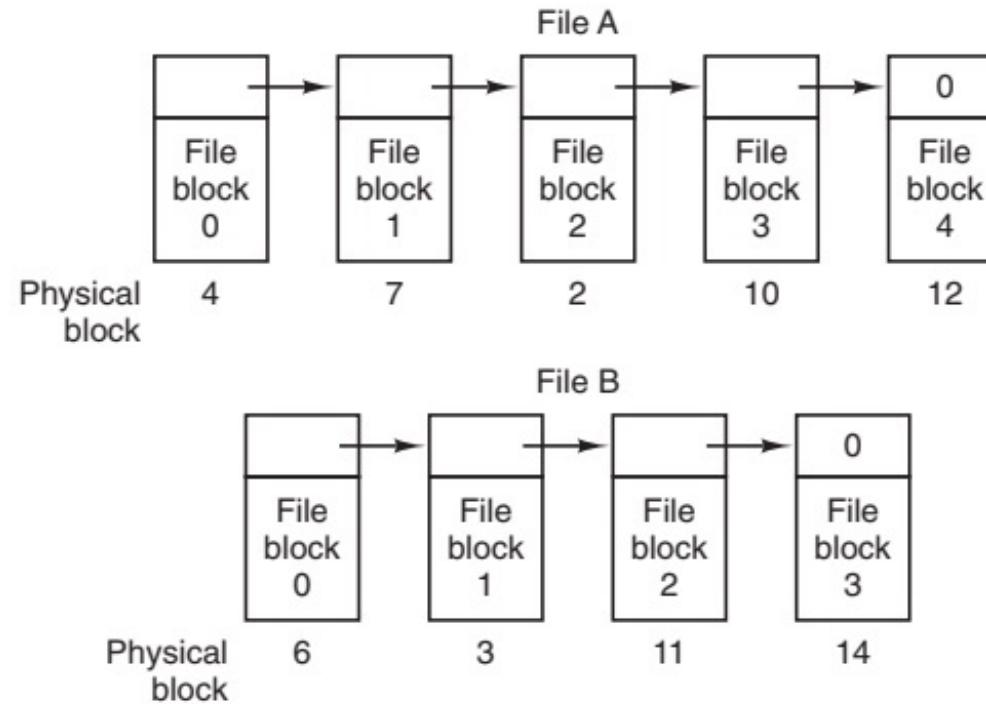
(b)

(a) Contiguous allocation of disk space for 7 files

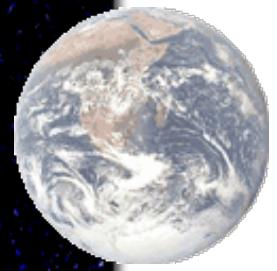
(b) The state of the disk after D and F have been removed



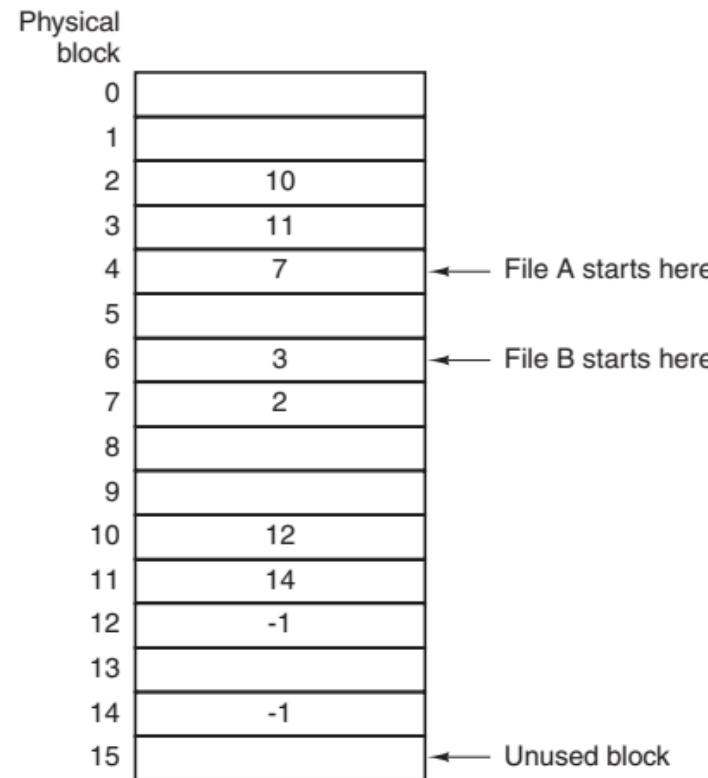
Linked-List Allocation



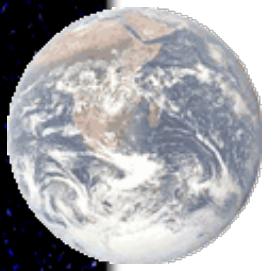
Storing a file as a linked list of disk blocks



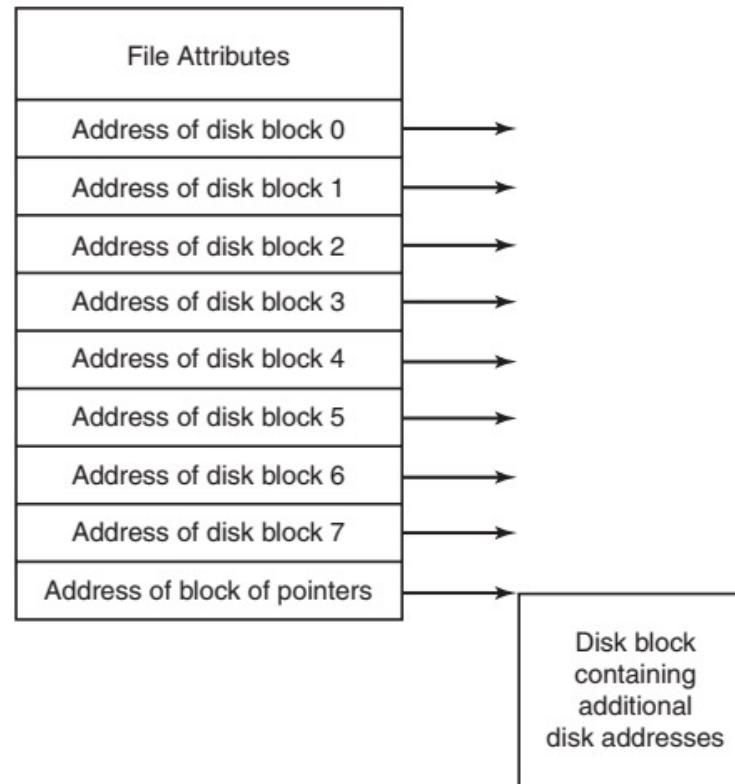
Link-List Allocation Using a Table in Memory

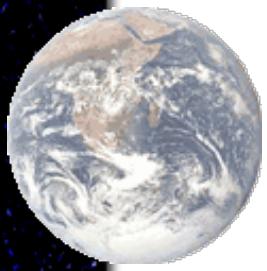


Linked list allocation using a file allocation table in memory

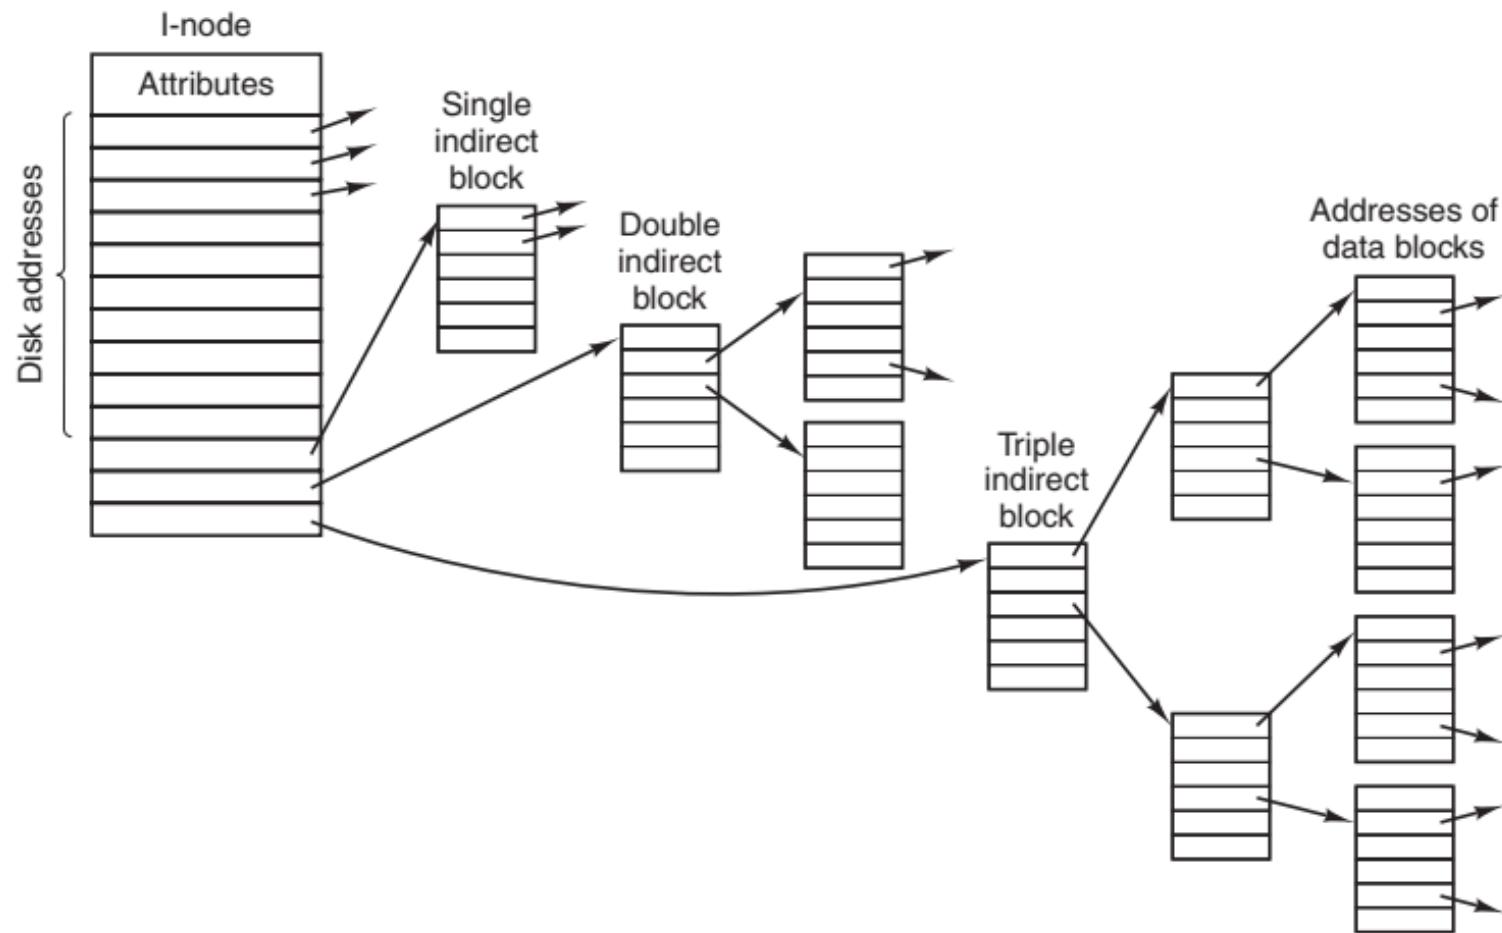


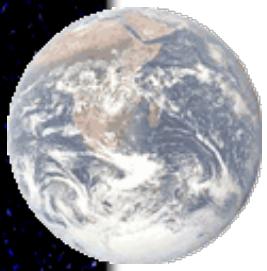
I-node





A UNIX I-node





The Steps in Lookin Up /usr/ast/mbox

Root directory

1	.
1	..
4	bin
7	dev
14	lib
9	etc
6	usr
8	tmp

Looking up
usr yields
i-node 6

I-node 6
is for /usr

Mode	size	times
132		

I-node 6
says that
/usr is in
block 132

Block 132
is /usr
directory

6	.
1	..
19	dick
30	erik
51	jim
26	ast
45	bal

/usr/ast
is i-node
26

I-node 26
is for
/usr/ast

Mode	size	times
406		

I-node 26
says that
/usr/ast is in
block 406

Block 406
is /usr/ast
directory

26	.
6	..
64	grants
92	books
60	mbox
81	minix
17	src

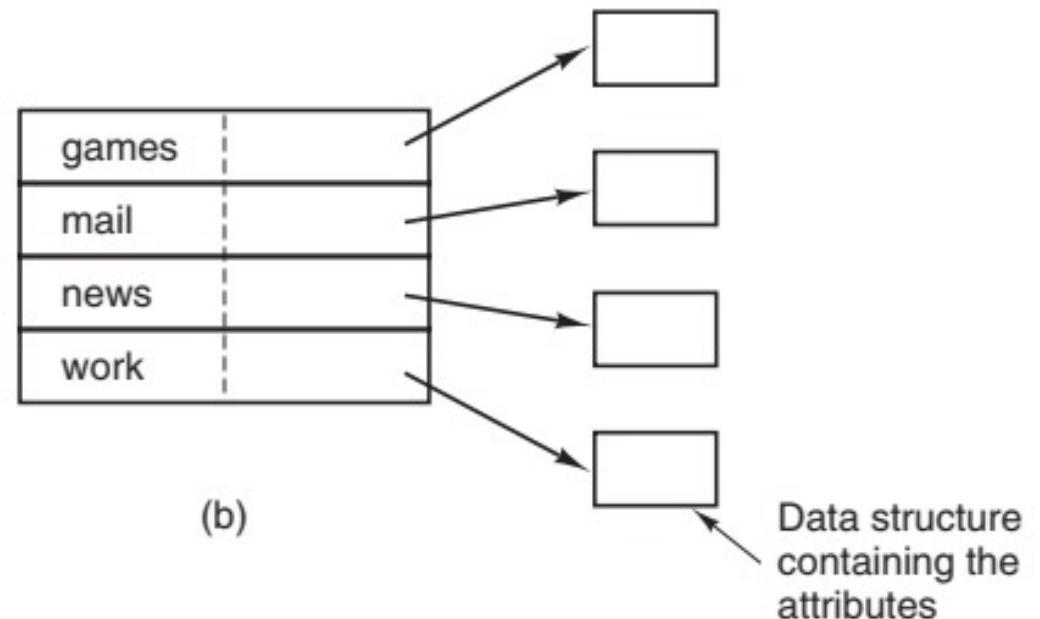
/usr/ast/mbox
is i-node
60



Implementing Directories

games	attributes
mail	attributes
news	attributes
work	attributes

(a)



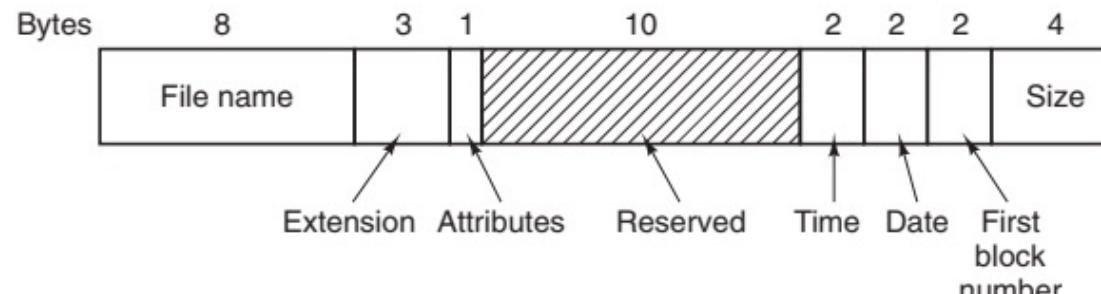
(b)

(a) A simple directory containing fixed-size entries with disk addresses and attributes in the directory entry

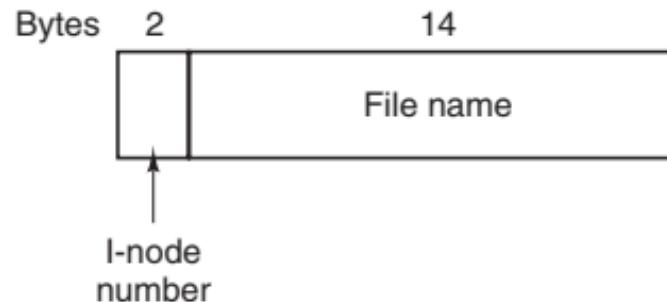
(b) A directory in which each entry just refers to an i-node



Examples of Directory Entry



The MS-DOS directory entry



A UNIX directory entry



Handling Long File Name

Entry for one file

File 1 entry length			
File 1 attributes			
p	r	o	j
e	c	t	-
b	u	d	g
e	t	<input checked="" type="checkbox"/>	
File 2 entry length			
File 2 attributes			
p	e	r	s
o	n	n	e
l	<input checked="" type="checkbox"/>		
File 3 entry length			
File 3 attributes			
f	o	o	<input checked="" type="checkbox"/>
⋮			

(a)

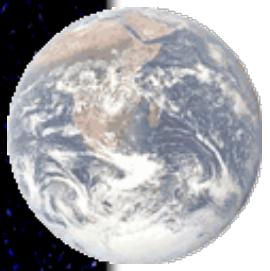
(a) In-line

Entry for one file

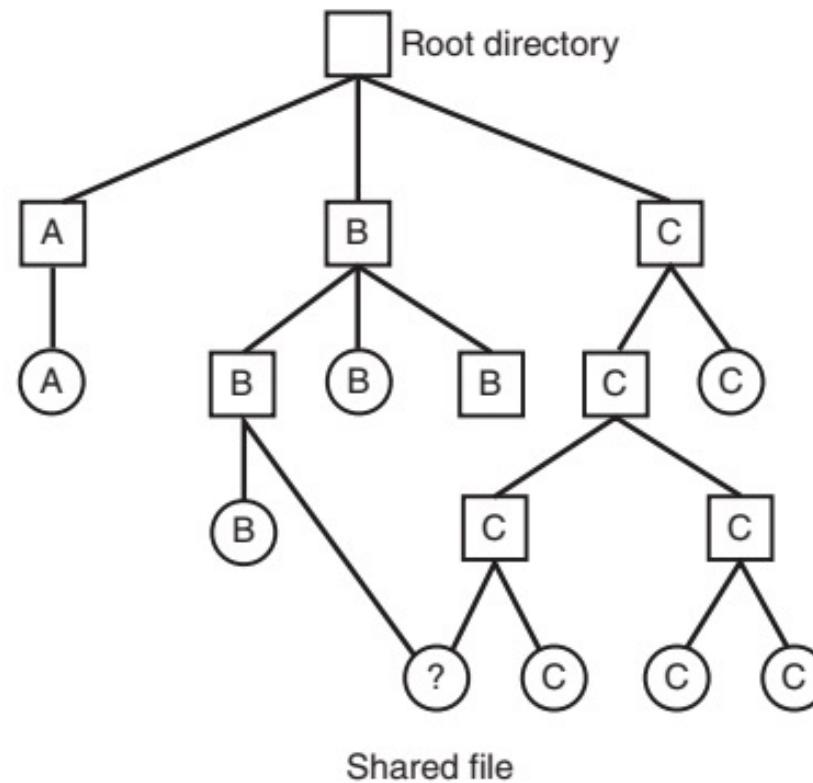
Pointer to file 1's name
File 1 attributes
Pointer to file 2's name
File 2 attributes
Pointer to file 3's name
File 3 attributes
p r o j
e c t -
b u d g
e t <input checked="" type="checkbox"/> p
e r s o
n n e l
<input checked="" type="checkbox"/> f o o
<input checked="" type="checkbox"/>

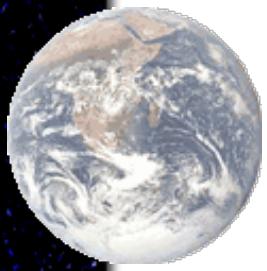
(b)

(b) In a heap



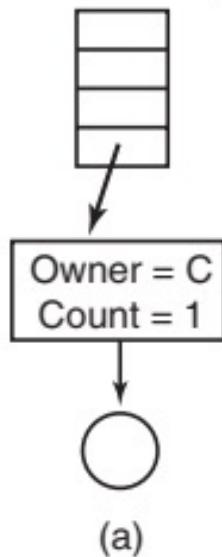
Shared Files



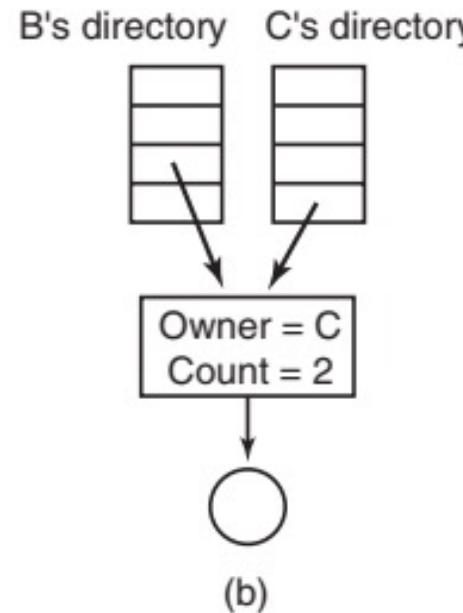


Shared Files

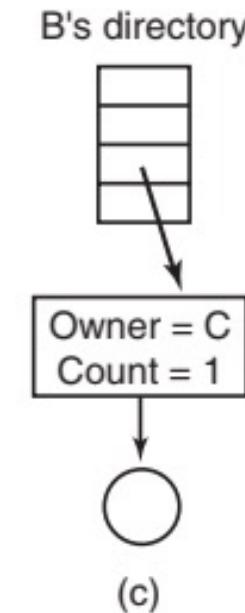
C's directory



B's directory



C's directory



(a) Situation prior to linking

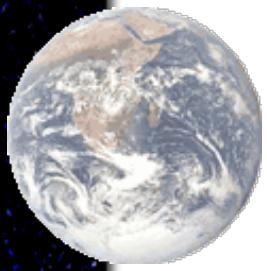
(b) After the link is created

(c) After the original owner removes the file



File System Management and Optimization

- Disk space management
- File-system backups
- File-system consistency
- File-system performance
 - Caching
 - Block read ahead
 - Reducing disk arm motion



Disk Space Management – Block Size

Length	VU 1984	VU 2005	Web
1	1.79	1.38	6.67
2	1.88	1.53	7.67
4	2.01	1.65	8.33
8	2.31	1.80	11.30
16	3.32	2.15	11.46
32	5.13	3.15	12.33
64	8.71	4.98	26.10
128	14.73	8.03	28.49
256	23.09	13.29	32.10
512	34.44	20.62	39.94
1 KB	48.05	30.91	47.82
2 KB	60.87	46.09	59.44
4 KB	75.31	59.13	70.64
8 KB	84.97	69.96	79.69

Length	VU 1984	VU 2005	Web
16 KB	92.53	78.92	86.79
32 KB	97.21	85.87	91.65
64 KB	99.18	90.84	94.80
128 KB	99.84	93.73	96.93
256 KB	99.96	96.12	98.48
512 KB	100.00	97.73	98.99
1 MB	100.00	98.87	99.62
2 MB	100.00	99.44	99.80
4 MB	100.00	99.71	99.87
8 MB	100.00	99.86	99.94
16 MB	100.00	99.94	99.97
32 MB	100.00	99.97	99.99
64 MB	100.00	99.99	99.99
128 MB	100.00	99.99	100.00

Percentage of files smaller than a given size (in bytes)

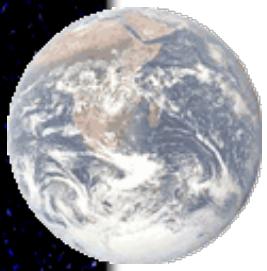


Disk Space Management – Block Size

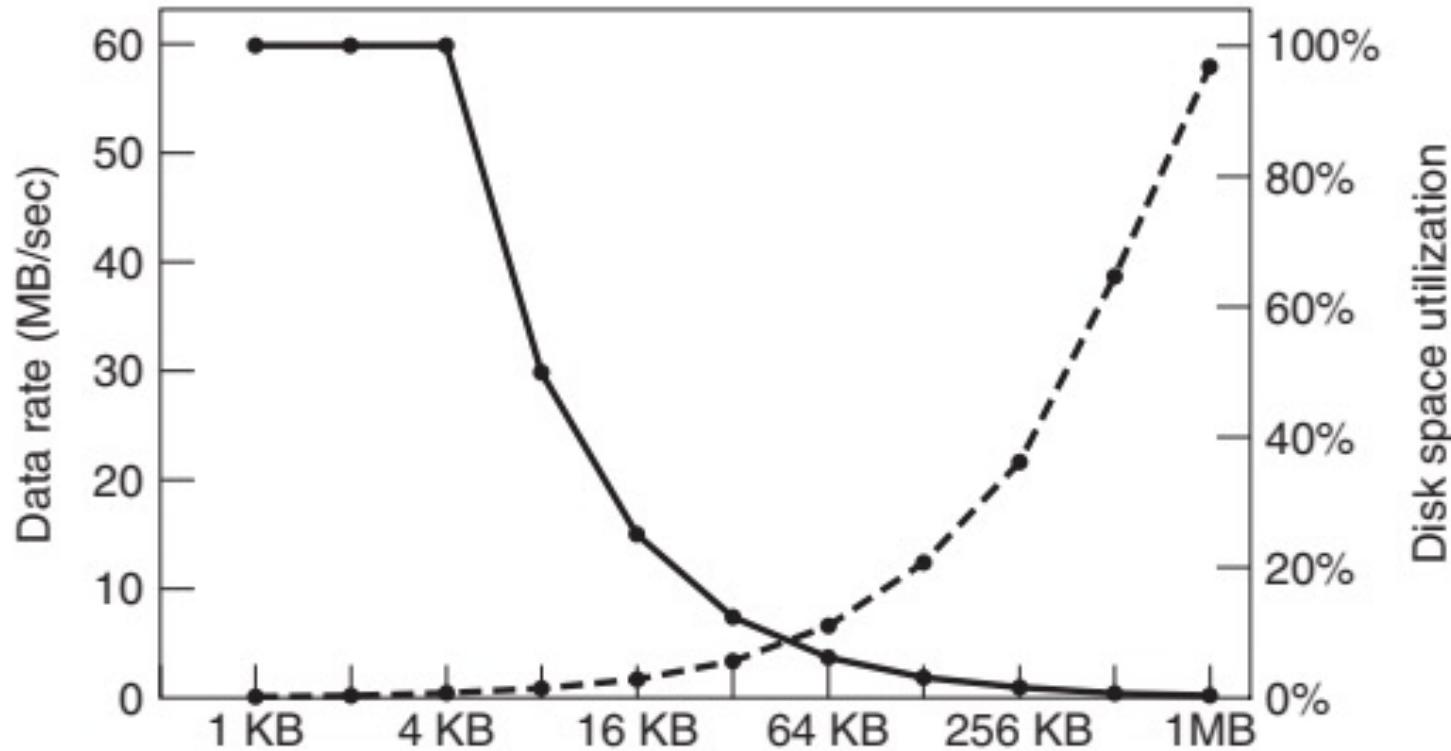
Length	VU 1984	VU 2005	Web
1	1.79	1.38	6.67
2	1.88	1.53	7.67
4	2.01	1.65	8.33
8	2.31	1.80	11.30
16	3.32	2.15	11.46
32	5.13	3.15	12.33
64	8.71	4.98	26.10
128	14.73	8.03	28.49
256	23.09	13.29	32.10
512	34.44	20.62	39.94
1 KB	48.05	30.91	47.82
2 KB	60.87	46.09	59.44
4 KB	75.31	59.13	70.64
8 KB	84.97	69.96	79.69

Length	VU 1984	VU 2005	Web
16 KB	92.53	78.92	86.79
32 KB	97.21	85.87	91.65
64 KB	99.18	90.84	94.80
128 KB	99.84	93.73	96.93
256 KB	99.96	96.12	98.48
512 KB	100.00	97.73	98.99
1 MB	100.00	98.87	99.62
2 MB	100.00	99.44	99.80
4 MB	100.00	99.71	99.87
8 MB	100.00	99.86	99.94
16 MB	100.00	99.94	99.97
32 MB	100.00	99.97	99.99
64 MB	100.00	99.99	99.99
128 MB	100.00	99.99	100.00

Percentage of files smaller than a given size (in bytes)



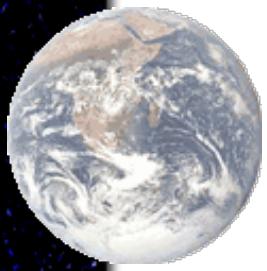
File System Backups



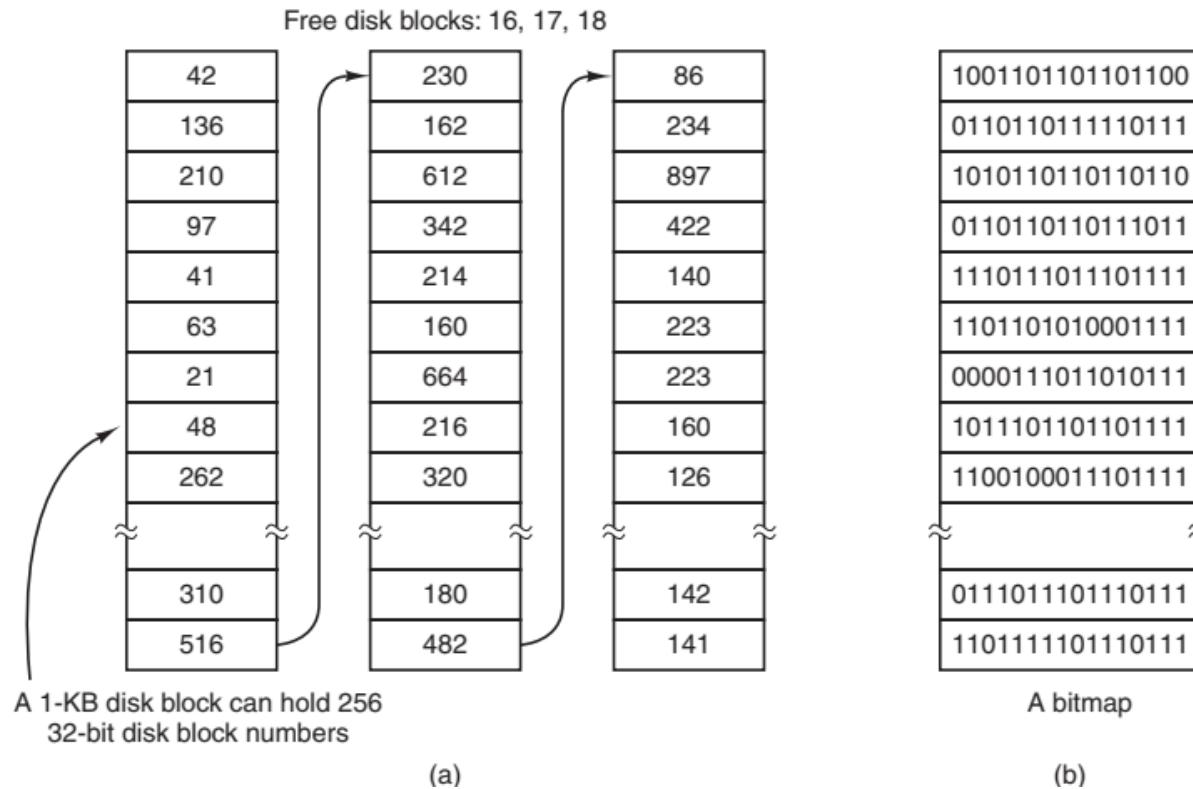
The dashed curve (left-hand scale) gives the data rate of a disk.

The solid curve (right-hand scale) gives the disk-space efficiency.

All files are 4 KB.

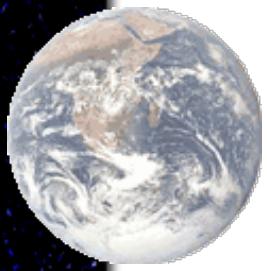


Keeping Track of Free Blocks

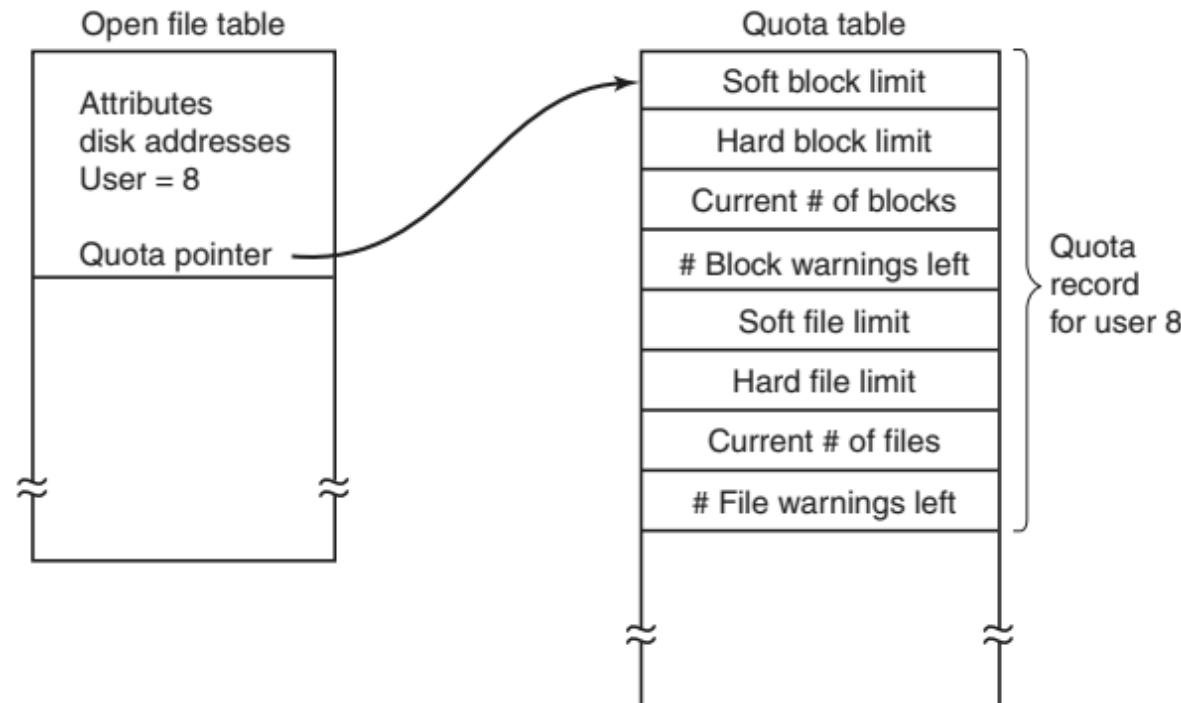


(a) Storing the free list on a linked list.

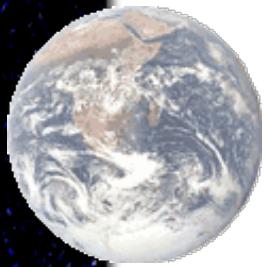
(b) A bitmap.



Disk Quotas



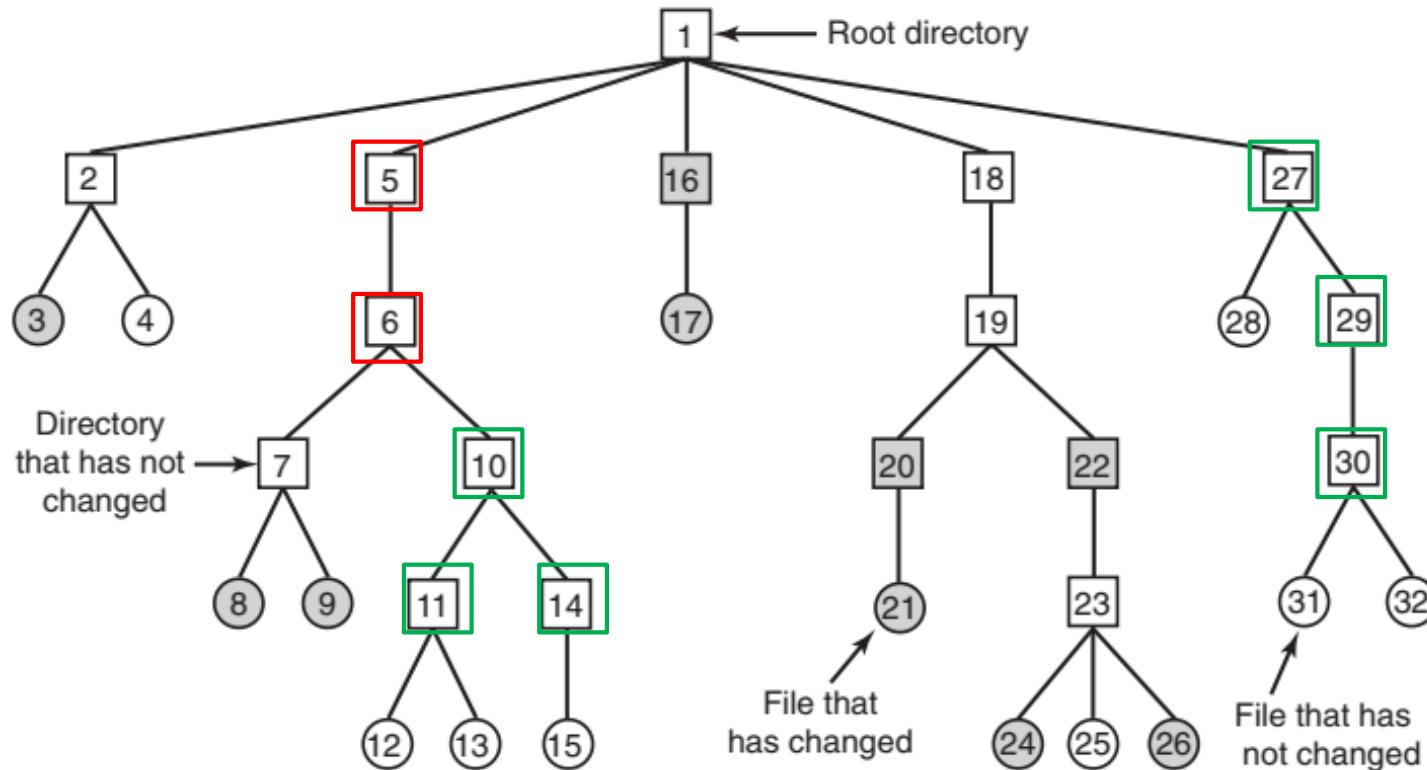
Quotas are kept track of on a per-user basis in a quota table.



File-System Backups

A file system to be dumped:

- The squares are directories and the circles are files.
- Each directory and file is labeled by its i-node number.



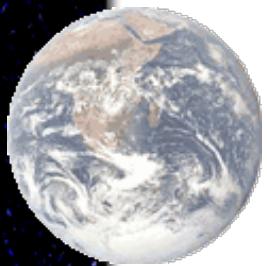
The shaded items have been modified since the last dump.



Bitmaps Used by the Logical Dumping Algorithm

(a)	<table border="1"><tr><td>1</td><td>2</td><td>3</td><td>4</td><td>5</td><td>6</td><td>7</td><td>8</td><td>9</td><td>10</td><td>11</td><td>12</td><td>13</td><td>14</td><td>15</td><td>16</td><td>17</td><td>18</td><td>19</td><td>20</td><td>21</td><td>22</td><td>23</td><td>24</td><td>25</td><td>26</td><td>27</td><td>28</td><td>29</td><td>30</td><td>31</td><td>32</td></tr></table>	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15	16	17	18	19	20	21	22	23	24	25	26	27	28	29	30	31	32
1	2	3	4	5	6	7	8	9	10	11	12	13	14	15	16	17	18	19	20	21	22	23	24	25	26	27	28	29	30	31	32		
(b)	<table border="1"><tr><td>1</td><td>2</td><td>3</td><td>4</td><td>5</td><td>6</td><td>7</td><td>8</td><td>9</td><td>10</td><td>11</td><td>12</td><td>13</td><td>14</td><td>15</td><td>16</td><td>17</td><td>18</td><td>19</td><td>20</td><td>21</td><td>22</td><td>23</td><td>24</td><td>25</td><td>26</td><td>27</td><td>28</td><td>29</td><td>30</td><td>31</td><td>32</td></tr></table>	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15	16	17	18	19	20	21	22	23	24	25	26	27	28	29	30	31	32
1	2	3	4	5	6	7	8	9	10	11	12	13	14	15	16	17	18	19	20	21	22	23	24	25	26	27	28	29	30	31	32		
(c)	<table border="1"><tr><td>1</td><td>2</td><td>3</td><td>4</td><td>5</td><td>6</td><td>7</td><td>8</td><td>9</td><td>10</td><td>11</td><td>12</td><td>13</td><td>14</td><td>15</td><td>16</td><td>17</td><td>18</td><td>19</td><td>20</td><td>21</td><td>22</td><td>23</td><td>24</td><td>25</td><td>26</td><td>27</td><td>28</td><td>29</td><td>30</td><td>31</td><td>32</td></tr></table>	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15	16	17	18	19	20	21	22	23	24	25	26	27	28	29	30	31	32
1	2	3	4	5	6	7	8	9	10	11	12	13	14	15	16	17	18	19	20	21	22	23	24	25	26	27	28	29	30	31	32		
(d)	<table border="1"><tr><td>1</td><td>2</td><td>3</td><td>4</td><td>5</td><td>6</td><td>7</td><td>8</td><td>9</td><td>10</td><td>11</td><td>12</td><td>13</td><td>14</td><td>15</td><td>16</td><td>17</td><td>18</td><td>19</td><td>20</td><td>21</td><td>22</td><td>23</td><td>24</td><td>25</td><td>26</td><td>27</td><td>28</td><td>29</td><td>30</td><td>31</td><td>32</td></tr></table>	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15	16	17	18	19	20	21	22	23	24	25	26	27	28	29	30	31	32
1	2	3	4	5	6	7	8	9	10	11	12	13	14	15	16	17	18	19	20	21	22	23	24	25	26	27	28	29	30	31	32		

- a) All modified files and all directories have been marked in the bitmap, as shown (by shading)
- (b) Unmarking any directories that have no modified files or directories in them or under them.
- (c) Dumping all the directories that are marked for dumping
- (d) Dumping all the files marked



File System Consistency

Block number	
0 1 2 3 4 5 6 7 8 9 10 11 12 13 14 15	
1 1 0 1 0 1 1 1 1 0 0 1 1 1 1 0 0	Blocks in use

Block number	
0 1 2 3 4 5 6 7 8 9 10 11 12 13 14 15	
0 0 1 0 1 0 0 0 0 1 1 0 0 0 1 1	Free blocks

(a)

Block number	
0 1 2 3 4 5 6 7 8 9 10 11 12 13 14 15	
1 1 0 1 0 1 1 1 1 0 0 1 1 1 1 0 0	Blocks in use

Block number	
0 1 2 3 4 5 6 7 8 9 10 11 12 13 14 15	
0 0 0 0 1 0 0 0 0 1 1 0 0 0 1 1	Free blocks

(b)

Block number	
0 1 2 3 4 5 6 7 8 9 10 11 12 13 14 15	
1 1 0 1 0 1 1 1 1 0 0 1 1 1 1 0 0	Blocks in use

Block number	
0 1 2 3 4 5 6 7 8 9 10 11 12 13 14 15	
0 0 1 0 2 0 0 0 0 1 1 0 0 0 1 1	Free blocks

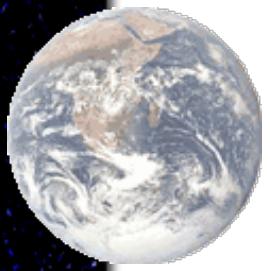
(c)

Block number	
0 1 2 3 4 5 6 7 8 9 10 11 12 13 14 15	
1 1 0 1 0 2 1 1 1 0 0 1 1 1 1 0 0	Blocks in use

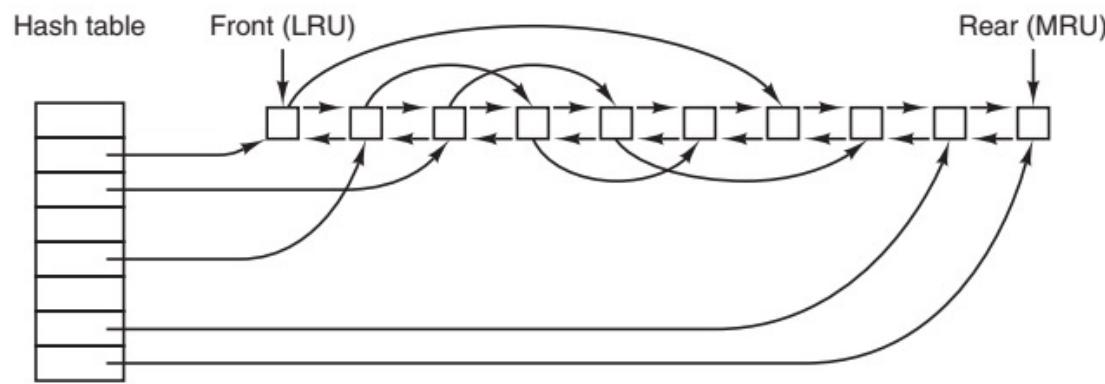
Block number	
0 1 2 3 4 5 6 7 8 9 10 11 12 13 14 15	
0 0 1 0 1 0 0 0 0 1 1 0 0 0 0 1 1	Free blocks

(d)

- (a) Consistent
- (b) Missing block
- (c) Duplicate block in free list
- (d) Duplicate data block



Caching

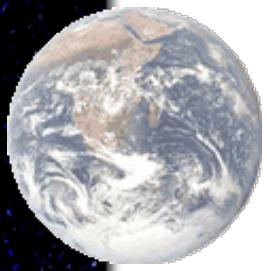


The buffer cache data structures

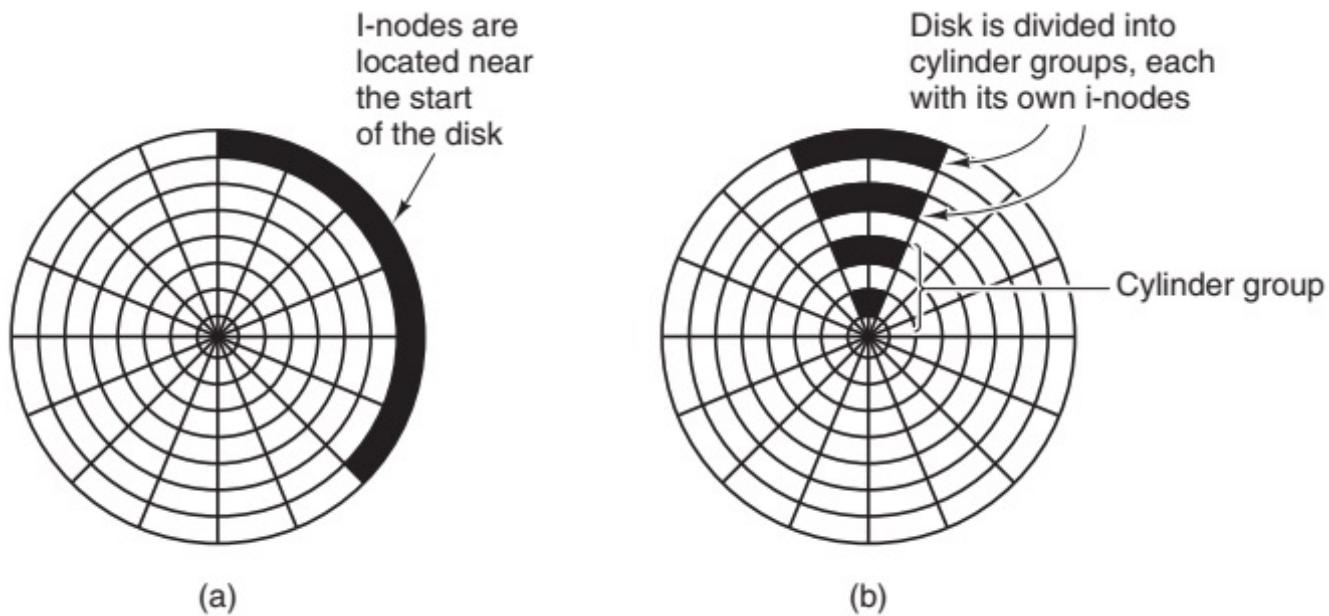


Block Read Ahead

- Get blocks into the cache before they are needed to increase the hit rate.
- For sequentially read file
 - Read block $k+1$ along with a read request of block k



Reducing Disk Arm Motion



(a) I-nodes placed at the start of the disk.

(b) Disk divided into cylinder groups, each with its own blocks and i-nodes.



References

1. Modern Operating Systems, 4th edition, Andrew S. Tanenbaum, Herbert Bos
2. Operating Systems, 3rd edition, H.M.Deitel, Pearson Education Limited: Longman.
3. Operating System Concepts, 10th edition, Abraham Silberschatz, Yale University, Peter Baer Galvin, Pluribus Networks, Greg Gagne, Westminster College, 10th edition. Wiley.
4. Operating Systems: Internals and Design Principles, 7th edition, William Stallings.